

**COVERING CODES CHECKED BY THE LEAN KERNEL:
 $K_7(9, 4) \leq 1134$, TEN MORE UPPER BOUNDS, $K_2(6, 1) = 12$,
AND $K_7(4, 2) = 19$ BY CERTIFIED SAT**

THIAGO PATZDORF

ABSTRACT. We prove in Lean 4, against Mathlib v4.34.1 and checked by the kernel with no `sorry` and no `native.decide`, upper bounds on the covering number $K_q(n, R)$ for eleven cells, each by an explicit code smaller than the corresponding entry in the tables we consulted. The largest improvement is $K_7(9, 4) \leq 1134$, against a best published bound of 1475; three more cells unchanged in the tables since 2011 come down to $K_7(10, 4) \leq 5607$ (from 6517), $K_5(11, 4) \leq 2875$ (from 3125) and $K_5(10, 5) \leq 162$ (from 175); the others are $K_7(8, 3) \leq 1887$, $K_5(10, 4) \leq 625$, $K_5(9, 3) \leq 1250$, $K_5(7, 2) \leq 500$, $K_5(9, 4) \leq 250$, $K_4(10, 4) \leq 192$ and $K_5(9, 5) \leq 50$. The codes are unions of cosets of a linear code plus a small patch, and we give two certificates: one walks digit prefixes and costs 12 CPU-hours for $(\mathbb{Z}/7\mathbb{Z})^9$; the other works on syndromes and costs minutes (the 5616-word code: 11 minutes on four cores). The base of the $K_7(9, 4)$ code was chosen from an enumeration of all 6362 equivalence classes of non-degenerate $[9, 3]_7$ codes. We also prove $K_2(6, 1) \geq 11$ by a short double-counting argument and the classical value $K_2(6, 1) = 12$ by a search whose soundness is a theorem. Finally we determine $K_7(4, 2) = 19$ (previously $17 \leq K_7(4, 2) \leq 19$): the upper bound and the reduction lemmas are kernel-checked, and the non-existence of an 18-word code is a SAT computation over 70 profile instances, refuted with LRAT proofs checked by two independent checkers and confirmed by a second, independent encoding; one of the 70 refutations is already checked in the kernel.

1. INTRODUCTION

The covering number $K_q(n, R)$ is the least cardinality of a code $C \subseteq A^n$, with $|A| = q$, such that every word of A^n lies at Hamming distance at most R from some word of C [1]. Upper bounds come from constructions and lower bounds from counting; the tables of Kéri [4] and the recent work of Gijswijt and Polak [3] and of Marosi [6] record the current state.

Upper bounds are certificates: a list of words and the claim that it covers. Checking the claim is a finite computation, but for q^n in the tens of millions it is done by programs that are rarely audited. Florath [2] built a Lean library of covering-code bounds whose derivations are replayed by the kernel; it formalizes the sphere-covering bound and records many cells. This note adds certificates that the kernel checks for eleven cells where the codes are new, a certificate that scales by working on syndromes, and a formal determination of a small classical value.

The proofs use Lean 4 [5] and Mathlib [7]. The source is <https://github.com/thiagopatzdorf/Matematica>, tag v0.7.0. A statement below is a claim only when it is the name of a declaration in that tree. The default target `CoveringLean` builds in seconds; the prefix certificates form the target `CoveringHeavy` and the syndrome certificates the target `CoveringSyn`.

2. DEFINITIONS

In the development `CoveringA2` a word is a function $\text{Fin } n \rightarrow \text{ZMod } q$, distance is Mathlib’s `hammingDist`, and

$$\text{Covers } R \ C : \iff \forall x, \exists c \in C, d(x, c) \leq R.$$

In `CoveringA6` the alphabet is `Fin q`, with the same distance and the same notion of covering, and

$$\text{IsK } q \ n \ R \ k : \iff (\exists C, |C| = k \wedge \text{Covers } R \ C) \wedge \forall C, |C| < k \rightarrow \neg \text{Covers } R \ C,$$

which is the equality $K_q(n, R) = k$.

3. UPPER BOUNDS FOR $K_7(9, 4)$

Theorem 1. *There is $C \subseteq (\mathbb{Z}/7\mathbb{Z})^9$ with $|C| = 1134$ that covers at radius 4. In Lean: `Syn.K7_9_4_le.1134_syn`, stated as*

$$\exists C : \text{Finset } (\text{Fin } 9 \rightarrow \mathbb{Z} \text{Mod } 7), \quad C.\text{card} = 1134 \wedge \text{Covers } 4 \ C.$$

The library proves the same statement for five codes of this cell, found in succession: 1351 words (`CoveringKernel.K7_9_4_le.1351_kernel`, and again as `Syn.K7_9_4_le.1351_syn`), 1285 words (`Syn.K7_9_4_le.1285_syn`), 1141 words (`Syn.K7_9_4_le.1141_syn`), 1137 words (`Syn.K7_9_4_le.1137_syn`) and 1134 words. Each code is a list of natural numbers $\sum_i w_i 7^i$ in strictly increasing order, and the same list is a text file in `data/codes/`; the canonical SHA-256 of the 1134-word code begins `a48e3a42`.

Structure. All five codes are three cosets $r_j + C_0$ of a linear $[9, 3]_7$ code C_0 (1029 words) together with a patch. Write C_0 as the kernel of $H = [I_6 \mid A]$; the three cosets cover exactly the points whose syndrome is within distance-4 reach of one of three syndromes, and the remaining syndromes are *orphans*, whose 343 points each must be covered by the patch. For the 1351-word code there are 27 orphans and 322 patch words; for 1285, 24 orphans and 256 patch words; for 1141, with A given by the rows 666, 065, 652, 643, 621, 615 and coset syndromes 0, 7708, 4191, there are only 6 orphans (2058 points) and 112 patch words; the 1137- and 1134-word codes keep the same base and need 108 and 105 patch words.

First certificate: digit prefixes. A Boolean function `go` walks the coordinates one at a time, keeping for each prefix the codewords still within reach and their remaining budget of mismatches; `go_sound` proves that `go` returning `true` implies that every word has a codeword at distance at most R . `covers_of_go` turns that into `Covers`, using that Mathlib’s distance agrees with the digit-wise count, and `cert_of_go` adds the cardinality, from a kernel check that every entry is below q^n and the list is strictly increasing. `go_split` fixes a prefix of coordinates in advance, so the computation becomes many leaf theorems, each closed by `decide +kernel`; for $K_7(9, 4) \leq 1351$ there are $7^4 = 2401$ leaves in 95 files, about 12 CPU-hours in total. The pipeline accepts the binary Hamming code of length 7 and rejects a copy of it with one word replaced; both are tests in the library.

Second certificate: syndromes. The second certificate uses the structure. `SynCheck` (core Lean only, no Mathlib) defines Boolean checks over a transversal of the q^{n-k} syndrome classes: each class gets a witness (coset, message) or is marked orphan; each orphan’s q^k points are checked against the list; and each coset word $r_j + mG$ is shown to be in the list. `SynBridge` (with Mathlib, of fixed size) proves the checks sound, transports coverage along the translation by C_0 , and concludes $\exists C, C.\text{card} = M \wedge \text{Covers } R \ C$ in the theorem `syn_cert`. The leaves are `decide +kernel` over blocks of 4096 classes and import only `SynCheck`. The 1351-word code costs about 8 CPU-minutes this way instead of 12 CPU-hours, and each leaf stays below 1.8 GB. Each certificate also contains an `example` in which a wrong witness is rejected by the kernel.

Remark. The best upper bound we found is $K_7(9, 4) \leq 1475$, by Marosi [6] (arXiv version 3, September 2026); Kéri’s tables [4] give 1843. The best lower bound we know is 264. That the bound is new is our reading of those sources on 3 October 2026, not something the kernel checks.

4. A LOWER BOUND WITHOUT SEARCH: $K_2(6, 1) \geq 11$

Lemma 2. *For distinct $a, c \in (\mathbb{F}_2)^6$, the balls of radius 1 about a and c meet in 0 or 2 words.*

Proof. At distance 1 the intersection is $\{a, c\}$; at distance 2 it is the two words between them; at distance at least 3 it is empty. In Lean the case $a = 0$ is a kernel decision over 64 words (`even_inter_zero`), and translation gives the general case (`inter_card_translate`). $\square$

Theorem 3. *Every code that covers $(\mathbb{F}_2)^6$ at radius 1 has at least 11 words (`CoveringA6.K_2.6_1_ge.11`).*

Proof. Let $K = |C|$ and let $m(y)$ be the number of codewords whose ball contains y , so $m(y) \geq 1$ and $\sum_y m(y) = 7K$. For each x put $S(x) = \sum_{y \in B(x)} m(y) = \sum_{c \in C} |B(x) \cap B(c)|$. Then $S(x) \geq 7$, and if $x \notin C$ every term is even by Lemma 2, so $S(x) \geq 8$. Summing $S(x) + [x \in C] \geq 8$ over all 64 words and using $\sum_x S(x) = 7 \sum_y m(y) = 49K$ gives $50K \geq 512$, hence $K \geq 11$. The counting is `excess_bound`, stated for any family of symmetric balls of size 7 with even pairwise intersections. $\square$

This replaces a hypothesis that the previous version of this note left open: the statement `no_cover_2.6.1.le10_uncond` is the earlier conditional one with the condition removed.

5. AN EXACT VALUE: $K_2(6, 1) = 12$

Theorem 4. *`SC.K_2.6.1_eq12 : IsK 2 6 1 12`.*

The value is classical [8]. The upper bound is an explicit 12-word code (`code12`, encodings 0, 1, 2, 15, 28, 23, 57, 58, 59, 39, 44, 52). For the lower bound, a search `chkN` on 64-bit masks takes a state (words still allowed, covered set, forbidden centres), picks the least uncovered word, and branches on the centres that could cover it, forbidding at each branch the siblings already tried. It prunes when the remaining words cannot cover what is left, at 7 words per centre. `chkN_sound` proves that a `true` answer means no completion exists. A translation argument lets one assume the word 63 is in the code. The search from that root is cut at depth 2 into 38 states, each refuted by `decide +kernel` in its own file. The Lean database of Florath [2] records $10 \leq K_2(6, 1) \leq 12$ at this cell.

6. AN EXACT VALUE BY CERTIFIED SAT: $K_7(4, 2) = 19$

Kéri's tables give $17 \leq K_7(4, 2) \leq 19$. The upper bound is the partition construction (the word 0000 together with an MDS $[4, 2, 3]_3$ code on each of the alphabets $\{1, 2, 3\}$ and $\{4, 5, 6\}$); in Lean it is `K742.K_7.4.2.le19`, decided by the kernel over the 2401 words. For the lower bound, two words of $\mathbb{Z}_7^4$ are at distance at most 2 exactly when they agree in at least two coordinates. A counting argument on a box of shortened words (Lemma 0) gives $K_6(3, 1) \geq 18$ and $K_7(3, 1) \geq 21$; with it, a fibre $\{c \in C : c_j = a\}$ of a code of radius 2 with at most 18 words has at least 2 words (Lemma 1, `K742.fiber_card_ge_two`). The multiset of fibre sizes per coordinate, up to permuting coordinates, is one of 70 *profiles* for $M = 18$ (`profiles18_card`), and each profile gives a CNF that is satisfiable whenever a code of that profile exists. Following the method Florath used for $K_8(4, 2) = 23$ [2], we refuted all 70 CNFs with CaDiCaL and LRAT proofs (3.0 GB in total), each checked by `lrat-check` and by an independent checker; the 15 profiles of $M = 17$ are refuted the same way. An independent encoding (one variable per word, coverage by radius-2 balls, a totalizer for the cardinality and only lexicographic symmetry breaking) refutes again all 15 profiles of $M = 17$ and all 70 of $M = 18$, with another solver; an adversarial review checked the lemmas, the list of profiles and the symmetry breaking, including by verifying that 1200 randomly relabelled codes satisfy the CNF of their profile and that four deliberately wrong symmetry-breaking constraints are caught. Together this gives $K_7(4, 2) = 19$.

Remark. In Lean the upper bound, Lemmas 0 and 1, the profile reduction, an LRAT checker with a soundness proof (`LratK`) and one of the 70 refutations (profile 64, `K742Sat.refut18.p64`) are kernel-checked, and `K742.K_7.4.2.eq.19_of` proves the equality from two named hypotheses: the bridge from a code to its CNF and the remaining refutations. Discharging both is work in progress; until then the lower bound is a certified computation, not a theorem of this note.

7. THE SPHERE-COVERING BOUND

The library also proves the sphere-covering bound $q^n \leq |C| \cdot V$ (`CoveringA2.sphere_covering`), its strict form when $V \nmid q^n$, and eight numerical instances `CoveringChain.SPH.K*_lb`, among them $K_7(9, 4) \geq 221$. This was the content of version 0.2 of this note. The bound had already been formalized by Florath [2], and all eight instances are below the best lower bounds in the literature (for $K_7(9, 4)$, 264), so they are not claimed as results.

8. TEN MORE UPPER BOUNDS

The same certificates prove ten more upper bounds (Table 1). Each declaration is stated as $\exists C, C.\text{card} = M \wedge \text{Covers } R \ C$ over $\text{Fin } n \rightarrow \text{ZMod } q$, and for each cell the list in Lean decodes to the code file in `data/codes/` with the same canonical SHA-256. For $q = 4$ our literature search was weaker than for $q \geq 5$, and $K_5(10, 4) \leq 625$ is a linear $[10, 4]_5$ code, which should still be compared with the tables of Davydov, Marcugini and Pambianco on linear covering codes.

cell	ours	previous	source	certificate
$K_7(9, 4)$	1134	1475	Marosi	syndromes (also 1137, 1141, 1285; 1351 by prefixes)
$K_7(10, 4)$	5607	6517	Kéri	syndromes (also 5616)
$K_5(11, 4)$	2875	3125	Kéri	syndromes
$K_5(10, 5)$	162	175	Kéri	syndromes
$K_7(8, 3)$	1887	2337	Kéri	syndromes (1893 by prefixes)
$K_5(10, 4)$	625	875	Kéri	prefixes
$K_5(9, 3)$	1250	1275	Kéri	prefixes
$K_5(7, 2)$	500	525	Kéri	prefixes
$K_5(9, 4)$	250	255	Kéri	prefixes
$K_4(10, 4)$	192	208	Kéri	prefixes
$K_5(9, 5)$	50	55	Kéri	prefixes

TABLE 1. Upper bounds proved in Lean and checked by the kernel.

The three new cells of Table 1 (`Syn.K7_10_4.le_5607_syn`, `Syn.K5_11_4.le_2875_syn`, `Syn.K5_10_5.le_162_syn`) have stood in Kéri's tables since at least 2004; Kéri's 2011 update and the captures we compared leave them unchanged, and we found no later publication of a smaller code (Section 11). The best lower bounds we know are 1007, 546 [3] and 41. Their canonical SHA-256 begin 7670a7bd, 2a700e2f and 74176af6.

9. HOW THE CODES WERE FOUND

The first codes came from the generator `lincov` of Marosi's public repository `Mapika/coldcase` [6] followed by greedy patching. The improvements came from treating the base as the main variable. For $K_7(9, 4)$ we enumerated all 6362 classes of non-degenerate $[9, 3]_7$ codes up to equivalence (the list is complete: a mass-formula check, $\sum_S |\text{PGL}(3, 7)|/|\text{Stab}(S)| = 31\,931\,793\,642$, equals the number of 9-point multisets of $\text{PG}(2, 7)$ that contain a frame) and computed, for the classes with the smallest uncovered set, the exact minimum number of orphan syndromes left by three cosets. The best, 6 orphans against 27 for the original base, gives the 1141-word code after a simulated-annealing

search on the residual points, and the 1137-word code after a large-neighbourhood search whose subproblems are solved exactly as integer programs. A ball of radius 4 contains at most 24 of the 2058 residual points, so this base needs at least $\lceil 2058/24 \rceil = 86$ patch words; we use 108, and the gap is open. An exact computation of the minimum over all 6362 classes, and over the 1375 classes of degenerate $[9, 3]_7$ codes, is under audit and is not a claim of this note. For the 1134-word code we restricted the patch to be invariant under translation along a line of C_0 (orbits of 7 words); an integer program whose variables are orbits of candidate patch words, solved to optimality by HiGHS for the two best lines (3.4 and 3.6 hours), shows that 105 words (15 orbits) is the minimum under that symmetry, so going lower needs a less symmetric patch or another base. The three new cells use the same scheme: $K_7(10, 4) \leq 5616$ is 16 cosets of a $[10, 3]_7$ code (5488 words, 3 orphan syndromes, 1029 points) plus 128 words, and the 5607-word code keeps the base with a patch of 119 words, found by a randomized greedy for unicost set cover over the 1029 residual points; $K_5(11, 4) \leq 2875$ is 23 cosets of an $[11, 3]_5$ code with no patch; $K_5(10, 5) \leq 162$ is one $[10, 3]_5$ code (125 words, 8 orphan syndromes, 1000 points) plus 37 words, where a linear-programming bound shows that this base needs at least 22 patch words. For $K_7(8, 3)$ the base of the 1893-word code is five cosets of an $[8, 3]_7$ code (1715 words) with two orphan syndromes (686 points); reoptimizing the patch gives 172 words, which is optimal for that base, since no ball of radius 3 contains more than four of the 686 points and $\lceil 686/4 \rceil = 172$. Every candidate was checked by programs independent of the one that found it before being certified in Lean; the codes, a structured description of each (base, cosets, patch, provenance) and an official checker are in the repository.

10. WHAT THE KERNEL ACCEPTED

On 2 October 2026, on a virtual machine of type `e2-highmem-8`, Lean 4.34.1 and Mathlib v4.34.1: `lake build` (default target, from a fresh clone) finished 8944 jobs in 4 min 16 s, peak 10.0 GB; `lake build CoveringHeavy` (the prefix certificates and $K_2(6, 1) = 12$) finished 9181 jobs in 1 h 54 min, about 9.3 CPU-hours, at most 9.4 GB per process; `lake build CoveringSyn` (the four syndrome certificates) finished: 8983 jobs in 10 min 35 s on 8 cores (37 CPU-minutes for all four), peak 6.9 GB per process. The fifth syndrome certificate, `CoveringLean.Syn_K1137` ($K_7(9, 4) \leq 1137$), built in 2 min 41 s on 8 cores (331 CPU-seconds), peak 6.7 GB. On 3 October 2026, in a container with 4 cores and 15 GB, the four certificates of version 0.6 built from scratch (their own modules; Mathlib, `SynCheck` and `SynBridge` from cache): `Syn_K1134` in 105 s, `Syn_K162` in 101 s, `Syn_K2875` in 364 s and `Syn_K5616` in 654 s of wall-clock time. On 4 October 2026, in the same container, `Syn_K5607` built in 1261 s of wall-clock time. For each of the four, three kinds of mutation were built and rejected by Lean: replacing one codeword, removing one codeword (both for the cardinality and for the coverage check), and corrupting one witness in the coverage table. `#print axioms` on Theorems 1, 3, 4 and on every declaration of Table 1 reports `propext`, `Classical.choice` and `Quot.sound` only.

11. LITERATURE SEARCH

For version 0.6 we collected, with a reproducible script, the works citing Kéri's tables, the book [1] and the main papers on q -ary covering codes, through OpenAlex, zbMATH Open, Crossref, OpenCitations and arXiv, down to two levels of citations (2103 works, 582 open-access PDFs read in full), together with Kéri's tables and their archived copies. No source we read gives a code smaller than ours in any of the four cells, nor a value of $K_7(4, 2)$. Papers behind a paywall were read only through abstracts and reviews, and a full-text index does not find numbers inside tables, so this is a search, not a proof of novelty.

12. WHAT IS NOT A THEOREM

Novelty is a reading of the literature, not a checked fact. The search programs are not verified; only their output is. For $K_7(4, 2) \geq 19$ the SAT encodings and the reduction to profiles are checked by the reviews and tests described in Section 6, not yet by the kernel. That the formal statements say what is intended is a human review: the reader should check the statements of Theorems 1, 3, 4 and the definitions of `Covers` and `IsK`. Two files in the tree, `A5b.Stress.lean` and `A6d.SearchHeavy.lean`, are not imported and are not claimed.

REFERENCES

- [1] G. Cohen, I. Honkala, S. Litsyn and A. Lobstein, *Covering Codes*, North-Holland, 1997.
- [2] A. Florath, *Formal foundations and proof-carrying certificates for q -ary covering codes in Lean 4*, arXiv:2606.09600, 2026.
- [3] D. Gijswijt and S. Polak, *Semidefinite lower bounds for covering codes*, arXiv:2504.01932, 2025.
- [4] G. Kéri, *Tables for bounds on covering codes*, <https://old.sztaki.hu/~keri/codes/>, 2011.
- [5] L. de Moura and S. Ullrich, *The Lean 4 theorem prover and programming language*, CADE 28, LNCS **12699**, Springer, 2021, 625–635.
- [6] M. Marosi, *New upper and lower bounds on covering codes $K_q(n, R)$ for alphabets of size $5 \leq q \leq 21$* , arXiv:2608.19872v3, 2 September 2026 (version 1, August 2026, had $K_7(9, 4) \leq 1743$).
- [7] The Mathlib Community, *Mathlib*, <https://github.com/leanprover-community/mathlib4>, release v4.34.1.
- [8] R. G. Stanton and J. G. Kalbfleisch, *Covering problems for dichotomized matchings*, Aequationes Math. **1** (1968), 94–103.

UNIVERSIDADE DE CAXIAS DO SUL, BRAZIL
 Email address: thiagosleman@gmail.com